

vNUIOA: Optimizing Non-Uniform I/O Access in Virtualized Environment for Multicore Systems

Jianmin Qian Jian Li Junshen Tan Yangde Li Zhiyuan Bo

ABSTRACT

This document is intended to serve as a sample for submissions to HPCA 2016. We provide some guidelines that authors should follow when submitting papers to the conference. In an effort to respect the efforts of reviewers and in the interest of fairness to all prospective authors, we request that all submissions follow the formatting and submission rules detailed below.

1. INTRODUCTION

Driven by the rapid growth of big data applications and cloud computing, most of today's applications are computing intensive or data intensive and these applications are running in the cloud to make full use of multi-core resources. To efficiently transfer data between high performance servers and these cloud applications become significantly important. Especially in the current high performance network environment, such inefficiency will cause huge performance degradation. However, the underlying topology of multi-core systems and virtualization pose great challenge in solving this problem. Several observations of performance degradation are reported because of hardware resources affinity [1, 2, 3]. Existing researches focusing on the memory locality [4, 5, 6] to reducing the remote access and virtual machine scheduling to minimize uncore penalty [7, 8, 9].

Each CPU socket accesses faster to the directly attached hardware devices than the remote ones, such as Non-Uniform memory Access (NUMA). However, directly attached devices are not limited to memory and CPU, I/O devices such as high speed Ethernet network adapter also connect to the CPU socket through the interconnect bus. When a thread access I/O devices remotely, latency will increased and inter-node maximum bandwidth decreased. This access asymmetry becomes main bottleneck of today's high performance computing. Especially in the current high performance virtualized environment such as Network Function Virtualization (NFV) [], I/O latency in the host NUMA machine can cause a huge performance degradation of the latency sensitive NFV workloads, which are running in the virtual machines (VMs). This significantly impacts the performance data center infrastructures.

Although previous studies presented several affinity models to optimizing NUMA system, these models are not suitable for current high performance virtualized environment for following reasons.

- Existing NUMA affinity models are 2-dimension and mainly focusing on the memory and virtual CPU locality, the consideration of high speed interconnections affinity such as I/O devices locality are poor. since almost the cloud space applications frequent transfer data with high speed I/O devices, current affinity model must take I/O devices into consideration.
- Many traditional NUMA performance model use the hop distance to introduce the access penalty between two devices. However, hop distance is not accurate for modeling, especially for the I/O performance modeling due to: First, the hop distance contain less details about the underlying topology such as the access latency, so it can not be used in the accurate performance measurement. Second, hop distance always be a constant in the system. When the performance of the system changing in real time, the hop distance can not reflect this change and cause the performance prediction inaccurate.
- Furthermore, previous work only optimizing the Non-Uniform I/O access by performance tuning. They just binding the virtual machine threads and its memory together to the I/O device attached nodes to ensure local I/O access. This static method is not always efficient because too many threads in one node will increase the CPU load and hurt the performance significantly when the system load is too high.

In this paper, we present a novel locality model, vNUIOA, which mainly optimizing IO locality in the high performance virtualized system and both taking thread to thread affinity in to consideration. We first install SR-IOV technology into our system to providing high performance I/O virtualization, then we implement a hardware monitor by change the performance monitor units (PMU) both in physical machine and virtual machines. Based on the PMU hardware resource usages such as IO requests, VM address distribution and CPU utilization, we analyzed We evaluated our model on Intel multi-core machine and conduct a series of experimental. Results show that

The contributions of this paper include:

This paper is organized as follows: Section 2 provides an overview of NUIOA architecture, I/O virtualization technology and motivation for this research. Section 3 characterizes performance in NUIOA virtualized systems. Sections 4 and 5 present our design and implementation. Section 6

evaluates our prototype and compares with existing methods. Section 7 discusses related work and Section 8 concludes this paper.

2. BACKGROUND AND MOTIVATION

In this section, we begin with a brief introduction to asymmetric access architecture in multicore systems. Then we show performance degradation due to asymmetric access in high performance network environment. Lastly, we talk about that virtualization also aggravate the performance degradation.

2.1 Asymmetric Access in Multicore Systems

Asymmetric Memory Access: Traditional multicore processor shared with one memory controller. With the number of cores per socket increases, one memory controller results in high memory controller contention and bandwidth competition. In order to improve the utilization of multicore system, each CPU socket has been designed with its own integrated memory controller (IMC) which have faster access to local memory bank than the remote ones. This Non-Uniform Memory Access (NUMA) characteristic fully take advantage of multicore system and high speed memory bandwidth.

Asymmetric I/O Access: Besides the asymmetric memory access, this architecture also result in asymmetric accesses of other devices. One of the most important one is I/O devices, because I/O devices are also directly connected to one or more nodes in the multicore system. local cores and devices have privilegiert access to the I/O device, remote ones must use interconnection to transfer data to the I/O devices, resulting in extra inter-node bandwidth occupation and access latency. This asymmetry I/O access is called Non-Uniform I/O Access (NUIOA)[]. AMD processors have exhibited NUIOA effects when AMD introduced the OPTERON processor in 2003 [], Intel processors have began to appear NUIOA characteristic when the advent of *sandy-bridge* architecture in 2011 []. Figure 1 shows a ubiquitous four-sockets NUIOA architecture based on the Intel Xeon E5 Family processors. Each socket consists of eight cores that share the last level cache (L3 cache), memory controller and physical memory bank. sockets link with each other via a high-speed, point-to-point interconnection such as Intel Quick Path Interconnect (QPI)[]. Network adapter is directly attached to the socket 0 with the Intel Data Direct I/O (DDIO) technology [], consequently, data transfer remotely to other sockets and increase access latency and bandwidth consumption.

2.2 Performance Degradation Due to Asymmetric Access

2.2.1 Moving from 10GbE to 40GbE

With the data transfer speed of current network adapter become higher and higher, such as Mellanox 40 and 56 Gigabit Ethernet adapter [] is being adopted broadly in the data center. The speed of these network adapter is higher than the memory and interconnection bandwidth, and result in shifting of system bottleneck from asymmetric memory access to asymmetric I/O access in the NUMA architecture. Finally, without consideration of the asymmetric I/O access in




Figure 1: Asymmetric I/O access architecture with four 8-cores Intel Xeon E5-4610 processors.

high speed networking environment, the performance of data center servers degrade significantly. To measure the Performance degradation due to asymmetric access, we conduct a set of experiments. First we evaluate the performance degradation of asymmetric I/O Access independently, then we combine asymmetric memory access and asymmetric I/O access and study the factors of performance degradation.

Table 1 shows the configuration details of our test hardware. We select Netperf [] benchmark provides network bandwidth testing between two host on a network.

Table 1: Configuration details of our test server.

Item	Configuration
CPU	Intel Xeon E5-4610 v2 2.3GHz (8 cores)
Memory	128G RAM DDR3 4 sockets, each with 32GB
QPI	7.2 GT/s, 2 links
Network Adapter	Mellanox ConnectX-3 Dual-Port 40 Gigabit Ethernet adapter
Hypervisor	KVM 2.0.0

2.2.2 Evaluation of Asymmetric I/O Access

To study the impact on performance of asymmetric I/O access independently, the VM vCPU and memory should be mapping to the same node to ensure local memory access. we simply run a number of VMs based on the asymmetric I/O access architecture by varying vCPU and memory mappings together among the multiple sockets. As shown in figure 2, a Four-sockets Intel Ivy Bridge NUMA architecture with ethernet adapter. Firstly, we banding vCPU of the VM to a specific core on the socket 0 and mapping all the memory of the VM to the memory bank attached to socket 0, then we change the VM vCPU banding to the core on the socket 2 and mapping memory of VM to the socket 2. Traditionally, these two kinds of mapping strategy ensure the VM vCPU thread is locally access to the VM memory, so the affinity between VM vCPU and memory can achieve best. However, we observe a huge I/O performance difference between this two kinds of VM vCPU and memory mapping strategy by using the Netperf benchmark.

Figure 3 gives the performance difference of these two kinds mapping strategies with Intel 10 Gigabit Ethernet adapter

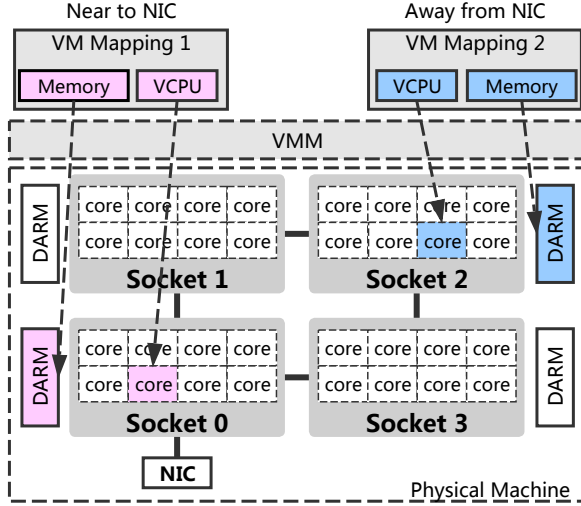


Figure 2: Different VM mapping strategies based on the asymmetric I/O Access architecture.

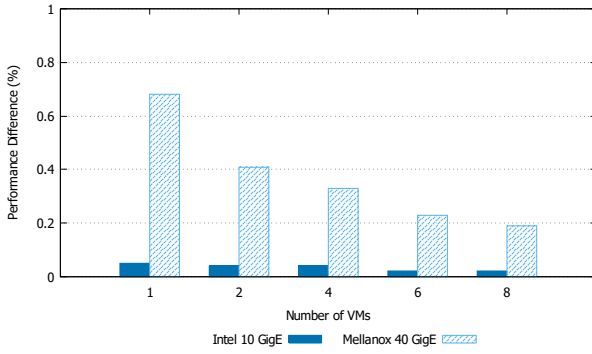


Figure 3: Performance degradation of different VM mapping strategies (VM mapping 2 relative to VM mapping 1) by using Intel 10 Gigabit Ethernet adapter and Mellanox 40 Gigabit Ethernet adapter.

and Mellanox 40 Gigabit Ethernet adapter. As for the Intel 10 Gigabit Ethernet adapter, we observe a little bit of throughput improvement (average 3%) because the speed of this network adapter can be fully taken advantage of. But for 40 Gigabit network adapter I/O throughput of first kind of mapping (VM mapping 1) is average 35.1% (max up to 68.1%) higher than the I/O throughput of second kind of mapping (VM mapping 2). The main reason of this huge performance difference is that when transfer data with 40 Gigabit network adapter, the speed of the network adapter is higher than QPI bandwidth. The second kind of mapping need to use two hops interconnection between sockets, this increase the QPI bandwidth consumption and memory access latency, finally, throughput will be decreased.

2.2.3 Evaluation of Asymmetric I/O Access and Memory Access

To study the factors of performance degradation in more complex scenarios, we consider the asymmetric I/O access and remote memory access simultaneously. Figure 4 shows three testing scenarios to characterize these two factors. In

the Figure 4(A), we bind all the VM vCPUs to the cores on socket 0 and mapping the memory respectively to socket 1, 2 and 3. This scenario ensure all the VM vCPUs in the I/O device attached socket and its memory are remote access. Similar in the Figure 4(B), we consolidate all the VM vCPUs on the socket 1 and mapping the memory respectively to socket 0, 2 and 3. This scenario is testing the performance of vCPU binding to the 1-hop distance socket and memory are remote access. Figure 4(C) show the VMs vCPU pinning to 2-hops distance socket and memory mapping to socket 0, 1 and 3. Not that we do not consider the VM vCPU and memory mapping to socket 3 because of that in our test bed socket 1 and socket 3 are symmetrical.

Figure 5 gives the evaluation results of these three scenarios.

2.3 Complexity in Virtualized Systems

Virtualization poses additional abstraction of the underlying topology and the mapping of VM vCPU and memory in the virtualized environment is more complexity. This complexity speedup the performance degradation seriously. Current VMMs tries hard to keep VM memory and vCPU as local as possible (like Xen and VMware ESXi), but mapping between virtual and machine resources must consider all kinds of factors as much as possible and different factors may disturb the distribution of VM vCPUs and memory. According to the white paper of VMware [], the main reason of chaotic mapping in the virtualized system can be listed as follows:

1. The heterogeneity of the virtual machine will impact the mapping strategies. Different kinds of application running on the VM changes the resources requirement. For example, memory intensive applications running in the virtual machine will result in the VM need larger size of memory, and these large amount of VM memory may distribute across multiple nodes.
2. The load balancing mechanism cause the VM memory and vCPU threads migration. Guest OS shutdown and reboot both trigger load balancing in the host OS, the prior mapping relationship will be changed. This will affect the affinity between vCPU, memory and network adapter. Finally, the system performance degrades.

The asymmetric I/O access degrade the I/O performance significantly in today's high speed network environment. Virtualization also poses a great challenge for optimizing the affinity among the vCPU, memory and network adapter. These motivate our work and in the next section, we further explore performance degradation in details by a set of experiments.

3. ANALYSIS OF OPTIMAL AFFINITY MODELING

In this section, we begin with the necessity and benefits of modeling the affinity in asymmetric access architecture. We then introduce a set of metrics that can be used to model the affinity in the asymmetric access architecture.

3.1 Necessity of optimal Affinity Modeling

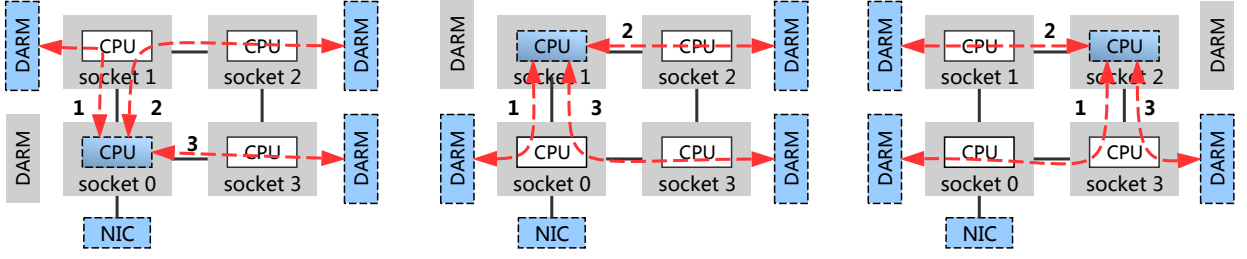


Figure 4: Three scenarios that we use to evaluate performance degradation of the asymmetric I/O access and remote memory access.

Poor Affinity Affect Performance: As we experimented and analyzed in the section 2, asymmetric access architecture can not guarantee local access, frequently remote access will cause the performance degradation in high performance networking environment. Nontransparent management in the virtualized system disturb the distribution of all kinds of resources. All these factors finally result in poor affinity among the VM vCPU, memory and network adapter, consequently, the inaccurate affinity model affects the networking performance, poor network performance result in the inefficient of data center servers. Therefore, optimal affinity model in the asymmetric access architecture is necessary.

Inaccurate of Previous Affinity Model: Historically, on the one hand, the previous affinity model in the asymmetric architecture mainly focus on optimizing the asymmetric memory access. Undoubtedly, these affinity models are no longer suitable for our asymmetric I/O access affinity modeling. On the other hand, previous affinity models are empirical [], they just characterize the penalty of asymmetric access by using the hop distances. Hop distances are typically recorded as a table in the BIOS Advanced Configuration and Power Interface (ACPI) to measure the distance between NUMA nodes. However, the hop distances lack details in some cases and are insufficient for affinity modeling in asymmetric access architecture.

3.2 Affinity Classes and Metrics

We now discuss different affinity classes in the asymmetric access architecture. To accurately determine the affinity among vCPU, memory and network adapter, we adopt more detailed metrics to describe the access behavior in the virtualized systems.

3.2.1 VCPU to Memory Affinity

VCPU to memory affinity refers to the traditional NUMA affinity between VM vCPU and its memory. Based on the previous work [], we standing in the point of view of the vCPU and expect the memory access locally as much as possible. The potential for NUMA affinity of vCPU is proportional to the amount of memory access to each NUMA node. That is, if a vCPU access most to the memory of a NUMA node, it has high affinity with this node. The NUMA affinity for vCPU can be calculated with equation 1, where A_m is the affinity between vCPU and each NUMA node, NC is a vector containing the accumulated number of memory access from VM vCPU to each node, and N is the

number of nodes.

$$A_m = \max(NC) / (\sum_{i=1}^N NC_i) \quad (1)$$

This affinity is minimal when the number of memory access from vCPU to each node is equal, in this case, the affinity equal to $1/N$. The affinity achieve maximum when all the memory access to one NUMA node, in this caes, NUMA affinity is 1.

3.2.2 VCPU to NIC Affinity

The vCPU to NIC affinity is defined as the physical access latency between NIC and the node which vCPU running on. Obviously, the node near to the NIC has best vCPU to NIC affinity. For latency sensitive virtual machines, it is better to binding all their vCPUs to the node near to the network adapter so as to achieve optimal performance. Pervious work [] has shown that CPI(cycles per instruction) can quantify the execution of a program. We use to this to measure affinity with the equation 2, where CPI_i denote the CPI of VM threads excuting on each NUMA node i (CPI_0 refers to CPI which excuting on the node near to NIC).

$$A_n = CPI_0 / (\sum_{i=1}^N CPI_i) \quad (2)$$

A_n achieve best when all the VM vCPU execute on the node near to NIC, in this case $A_n = 1$. When VM do not execute on the node 0, the A_n achieve minimum and equal to 0.

3.2.3 NIC to Memory Affinity

The NIC to memory affinity is the affinity between NIC ring buffer and VM memory buffer. When a NIC receives incoming packets, it copies the data packets into the buffers using DMA. Obviously, if the DMA buffer is allocated in the memory attached to the node which near to NIC, the access latency of NIC to memory is lowest. Otherwise, NIC access to memory buffer of the remote node by using the QPI link, in this caes, the QPI bandwidth consume quickly. We define the affinity between NIC buffer and memory by using the equation 3, where VMD is the VM memory distribution table, VMD_0 is memory of VM distribute on node 0 (which is near to the NIC), and $\sum_{i=1}^N VMD_i$ is the total number of VM

memory.

$$A_b = VMD_0 / (\sum_{i=1}^N VMD_i) \quad (3)$$

This affinity achieve best ($A_b = 1$) when all the VM memory are allocated on the node near to the NIC, the minimal value of A_b is 0 when the VM do not allocate any memory on the node near to the NIC.

3.3 Optimal Affinity Model Function

After calculate these three classes of affinity, we now discuss the optimal affinity model. In consideration of complexity in the virtualized environment, we first simply defined our optimal affinity model function as follows:

$$f(A_m, A_n, A_b) \quad (4)$$

As for I/O intensive virtual machines, data interacting with the host machine is frequently. When the NIC receives a large amount of packets, interrupts raising from the NIC also handled by the vCPU of the VM, packets were copied to the memory buffers by using DMA. In this case, the asymmetric I/O accesses are main factors which influencing the performance, our model function focus on optimizing the affinity A_n and A_b . materialization

$$f = (\alpha \times (A_n + A_b) + (1 - \alpha) \times A_m) \quad (5)$$

As for CPU intensive virtual machine, the vCPU threads frequently access memory to read and write data. In this case, the main factor resulting in the performance degradation is asymmetric memory access, our model function should focus on optimizing the affinity A_m .

$$f = (\beta \times A_m + (1 - \beta) \times (A_n + A_b)) \quad (6)$$

but not the maximum affinity will achieve best performance, because maximum data locality will increase cache miss rate and affects performance.

4. DESIGN

In this section, we describe the detail design of our vNUIOA, a platform independent runtime system that optimizing the overheads of asymmetric access architecture. We enabling three parts for incorporating NUMA and NUIOA overheads awareness with hypervisor: namely, Runtime performance analyzer, Affinity aware algorithm and Dynamic load balance scheduler.

4.1 Overview

An overview of vNUIOA is shown in Figure 5. It consists of the following three parts: the **Runtime Performance Analyzer** modules collects and analyzes the information of guest and host os since the system booting, then the module classify and store these information of each VM. After these information is collected, we design the asymmetric access overheads awareness algorithm based on the optimal affinity model which describe in the section 3.

4.2 Runtime Performance Analyzer

Runtime information collector mainly collects the information we used in the optimal affinity modeling. To periodically

collect the performance information of virtual machines and host server, it is imperative to analyze memory and I/O access activities in the virtualized system. To end this, we use the Linux performance monitoring tool `perfmon` [] to detect each virtual machine's memory and I/O access activities, then we store these information in the VMinfo table. Each table entry contains information about each NUMA node access to VM pages vector (NodeAcc Vector), VM memory distribution table (VMD table) and VM I/O request times vector (VIOR Vector). The detail of these three information listed as follows:

Table 2: Definitions of collected information.

Per VM statistics	
NC_n	NC_n records the accumulated number of memory access from VM vCPU to each NUMA node.
VMD_n	VMD_n vector describes the distribution of VM memory to node n
CPI_n	CPI_n vector describes the execution of VM vCPU on node n
NIOR	NIOR records the number of I/O requests per second for VM .
Physical machine statistics	
L_i	CPU load of core i ($L_i = \frac{CPUtime_i}{Elapsedtime_i}$).
DL_{mn}	Access latency between each NUMA node.

System load balance is also important for performance of physical server. As shown in the section 3.2.1, physical machine system load imbalance will cause a huge performance degradation. So in our design, we must take the system load into consideration. In this paper, we mainly consider average CPU load of each NUMA node, the per-node CPU load L_n can be calculated as follows:

$$L_n = (\sum_{i=1}^M L_i) / M \quad (7)$$

where M is the number of cores per NUMA node. we test that when the system load per node exceed the 80% of the max load, the system performance begin to goes down. Based on this experimental value, if the system load per node up to 80%, the load of the node is high.

Because of the hop-distance result in the previous NUMA affinity modeling inaccurate, in this work, we use a more accurate system information for NUIOA affinity modeling. We measure the average access latency between NIC and every NUMA node, then we record it in a DL vector ($DL[n]$). Based on the DL vector, the data transfer latency between NIC and nodes can be accurate measured.

4.3 Affinity aware algorithm

Step 1: we first identify which type a virtual machine belong to, because for different type VM, we adopt different optimal function to manage affinity. We classify the VMs into I/O intensive and I/O non intensive by looking the V-IOR vector (see the Table 2), the I/O request threshold is to be determined experimentally. We found that the number of I/O request for VM with running CPU intensive task is less

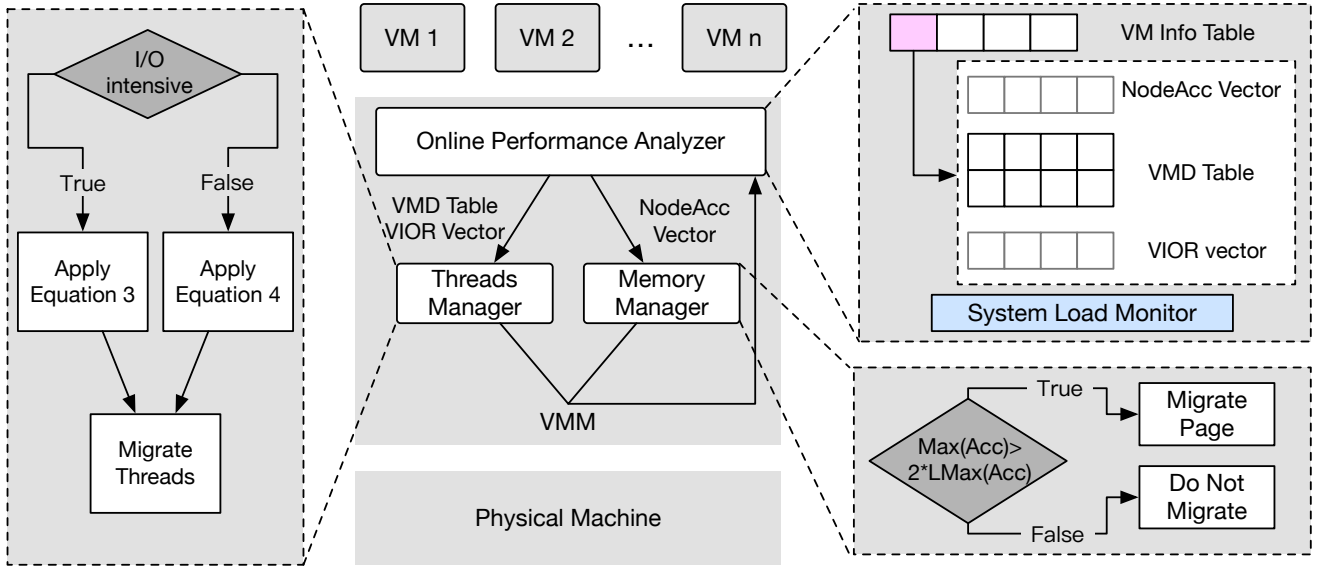


Figure 5: Architecture of vNUIOA .

than 100 per second and for I/O intensive task the number is more than 1000 per second, so we choose a threshold value of 500 times per second in consideration of practical application attributes.

Step 2: calculating the optimal affinity for each VM. For different type of virtual machine, we assign different weight to the three affinity A_m, A_n, A_b , in this way, we can figure out optimal affinity function. I/O intensive VMs always have frequent I/O operation, so the affinity A_n, A_b are more important in the optimal function, and as a result, the weight of the A_n, A_b is dominant. CPU intensive VMs always read and write data with memory, in this case, the proportion of affinity A_m is higher in the optimal function and we should assign a proper weight.

Step 3: According to the calculating results in the step 2, we make the decision of memory migration and vCPU mapping. we only migrate memory when maximum number of remote node memory access is double than the number of local memory access ($\max(NC) > 2 * NC[0]$), because frequent memory migration also bring high overheads. We test several value and found that when select 2 times than the number of local access, we can achieve max benefits. The vCPU mapping also negotiate with the load balance mechanism in the system, the load per cpu core can be calculated by $L_i = \frac{CPUtime_i}{Elapsedtime_i}$. If the average CPU load per node is higher than 80%, we should not map any vCPU to this node and select another node for vCPU remapping. We remapping the VM m vCPUs from the original node to node k instead of the node p is beneficial only if the average access latency to node k is less than average access latency to node p. remember L_k and L_p both less than 80%.

Step 4:

5. IMPLEMENTATION

We implemented the vNUIOA prototype in the KVM virtualized platform. The online performance analyzer, the thread-

Algorithm 1 Affinity aware Scheduling Algorithm

```

1: Variables:  $NIOR, L_n, NC_n$ 
2: for each VM m do
3:   /*Periodic look at  $NIOR$  vector*/
4:   if  $NIOR \geq 500$  then
5:      $f_{IO} = \alpha \times (A_n + A_b) + (1 - \alpha) \times A_m$ ;
6:     if  $L_n \leq 0.8$  then
7:       map vCPUs of VM m to node n
8:     else
9:       if  $DL_{nk} \leq DL_{np}$  then
10:        map VM vCPUs to the node k
11:       end if
12:     end if
13:   else
14:      $f_{NIO} = \beta \times A_m + (1 - \beta) \times (A_n + A_b)$ ;
15:     if  $\max(NC) \geq 2 \times NC[0]$  then
16:       migrating memory of VM m from to node n.
17:     else
18:       Do not migrate memory.
19:     end if
20:   end if
21: end for

```

s migration manager and the memory migration manager are implemented as individual daemons in the KVM system. We describes the implementation details of these three modules.

We use the KVM 2.0.0 as virtual machine monitor with the Intel VT-d and VT-x technology enable. Each VM run the Ubuntu 14.04 LTS system and configured with one vCPUs and 1GB memory. In order to achieve high performance I/O virtualization, we enable Single Root I/O virtualization (SR-IOV)[] technology and assigned a virtual function to each VM. SR-IOV [] proposes a set of hardware enhancements for PCIe device and enable to create multiple "lightweight" PCIe functions, known as virtual Functions (VF-

s) that still contain the major device resources. In this work, we mainly consider SR-IOV because With the hardware enhancement, SRIOV remove the major VMM intervention and achieve I/O virtualization without losing performance.

5.1 Techniques for Information Collection

Using page faults measure memory accesses: The OS has a methods to measure if the pages has been accessed by the page fault, While the page is not set as present in the page table, the OS is noticed about the faulting virtual address, as well as the thread ID that caused the fault. By tracking page faults, it possible to detect which vCPU threads are accessing the memory page. The performance analyzer accesses the hardware performance counters via Linux performance monitor tool *perfmon*. We using

Counting I/O requests per second: As for the number of I/O request per second, we count the execution number of *ixgbevf_clean_tx_irq* or *ixgbevf_clean_rx_irq* function for every second. Because every time a VM handles a packet, the SR-IOV vf driver invokes these two functions.

Migrating pages via system call: After determine memory migration, we use the *numa_migrate_pages()* system call to migrate all pages of the specified process from one node to another node. This system call supported by **libnuma** in the **numactl** package.

5.2

5.3

6. EVALUATION

In this section we present experimental evaluation of the proposed vNUIOA system, we compare the performance of vNUIOA with KVM's default scheduler and a hand-optimized binding strategy. Then we study the stability of our vNUIOA system. Finally, we characterize vNUIOA runtime overhead.

6.1 Experimental Environment

All experiments are conducted in a 4-sockets machine with Mellanox ConnectX-3 40 GigE adapter placed on the PCIe Gen3 x16 slots which directly connect to the socket 0. Each socket contains a Intel Xeon E5 4610 v2(Ivy Bridge architecture) processor and 32GB DDR3 memory bank. Each processor with 8 cores running at 2.3GHz and max turbo frequency can up to 2.7GHz. Each core has 32KB L1 data cache and 32KB L1 instruction cache, 256KB L2 cache and 16MB last level cache (LLC). Each socket equipped with 2 ways QPI with the communication speed 7.2GT/s. In our experiments, we enable the Intel Hyper-threading [(16 logical cores per socket). Therefor, we have total 64 cores and 128GB memory. we also configure the storage with 300G disk, the experiment configuration details are listed in the table 1.

We select a set of I/O and memory performance testing benchmarks,listed as table 2. The Netperf [] benchmark provides network bandwidth testing between two host on a network. Memcached [] is an in-memory key-value store for small chunks of arbitrary data to improve the performance of cloud applications.

6.2 Improvement of Performance

6.3

7. RELATED WORK

8. CONCLUSION

9. PRIOR/CONCURRENT SUBMISSIONS

10. REFERENCES